

Health_Care_Project

December 16, 2022

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import sklearn as sk
import warnings
warnings.filterwarnings("ignore")
%matplotlib inline
```

1 Perform preliminary data inspection and report the findings on the structure of the data, missing values, duplicates, etc.

```
[2]: df = pd.read_excel("HealthCare_dataset.xlsx")
```

```
[3]: df.head()
```

```
[3]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	\
0	63	1	3	145	233	1	0	150	0	2.3	0	
1	37	1	2	130	250	0	1	187	0	3.5	0	
2	41	0	1	130	204	0	0	172	0	1.4	2	
3	56	1	1	120	236	0	1	178	0	0.8	2	
4	57	0	0	120	354	0	1	163	1	0.6	2	

	ca	thal	target
0	0	1	1
1	0	2	1
2	0	2	1
3	0	2	1
4	0	2	1

```
[4]: df.tail()
```

```
[4]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	\
298	57	0	0	140	241	0	1	123	1	0.2	
299	45	1	3	110	264	0	1	132	0	1.2	

300	68	1	0	144	193	1	1	141	0	3.4
301	57	1	0	130	131	0	1	115	1	1.2
302	57	0	1	130	236	0	0	174	0	0.0

	slope	ca	thal	target
298	1	0	3	0
299	1	0	3	0
300	1	2	3	0
301	1	1	3	0
302	1	1	2	0

```
[5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         303 non-null    int64
1   sex         303 non-null    int64
2   cp          303 non-null    int64
3   trestbps    303 non-null    int64
4   chol        303 non-null    int64
5   fbs         303 non-null    int64
6   restecg     303 non-null    int64
7   thalach     303 non-null    int64
8   exang       303 non-null    int64
9   oldpeak     303 non-null    float64
10  slope       303 non-null    int64
11  ca          303 non-null    int64
12  thal        303 non-null    int64
13  target      303 non-null    int64
dtypes: float64(1), int64(13)
memory usage: 33.3 KB
```

2 Based on these findings, remove duplicates (if any) and treat missing values using an appropriate strategy

```
[6]: df.isnull().sum()
```

```
[6]: age         0
sex           0
cp            0
trestbps      0
chol          0
```

```

fbs          0
restecg      0
thalach      0
exang        0
oldpeak      0
slope        0
ca           0
thal         0
target       0
dtype: int64

```

```
[7]: df.duplicated().sum()
```

```
[7]: 1
```

```
[8]: df[df.duplicated()]
```

```

[8]:      age  sex  cp  trestbps  chol  fbs  restecg  thalach  exang  oldpeak  \
164   38    1   2      138    175    0          1      173     0         0.0

      slope  ca  thal  target
164      2   4     2        1

```

```
[9]: data = df.drop([164],axis = 0)
```

```
[10]: data.duplicated().sum()
```

```
[10]: 0
```

3 Get a preliminary statistical summary of the data and explore the measures of central tendencies and spread of the data

```
[11]: data.describe()
```

```

[11]:      age      sex      cp      trestbps      chol      fbs  \
count  302.00000  302.00000  302.00000  302.00000  302.00000  302.00000
mean    54.42053    0.682119  0.963576  131.602649  246.500000  0.149007
std      9.04797    0.466426  1.032044   17.563394   51.753489  0.356686
min     29.00000    0.000000  0.000000   94.000000  126.000000  0.000000
25%     48.00000    0.000000  0.000000  120.000000  211.000000  0.000000
50%     55.50000    1.000000  1.000000  130.000000  240.500000  0.000000
75%     61.00000    1.000000  2.000000  140.000000  274.750000  0.000000
max     77.00000    1.000000  3.000000  200.000000  564.000000  1.000000

      restecg      thalach      exang      oldpeak      slope      ca  \

```

count	302.000000	302.000000	302.000000	302.000000	302.000000	302.000000
mean	0.526490	149.569536	0.327815	1.043046	1.397351	0.718543
std	0.526027	22.903527	0.470196	1.161452	0.616274	1.006748
min	0.000000	71.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	133.250000	0.000000	0.000000	1.000000	0.000000
50%	1.000000	152.500000	0.000000	0.800000	1.000000	0.000000
75%	1.000000	166.000000	1.000000	1.600000	2.000000	1.000000
max	2.000000	202.000000	1.000000	6.200000	2.000000	4.000000

	thal	target
count	302.000000	302.000000
mean	2.314570	0.543046
std	0.613026	0.498970
min	0.000000	0.000000
25%	2.000000	0.000000
50%	2.000000	1.000000
75%	3.000000	1.000000
max	3.000000	1.000000

```
[12]: data.skew()
```

```
[12]: age          -0.203743
sex            -0.786120
cp              0.493022
trestbps       0.716541
chol           1.147332
fbs            1.981201
restecg        0.169467
thalach        -0.532671
exang          0.737281
oldpeak        1.266173
slope          -0.503247
ca             1.295738
thal           -0.481232
target         -0.173691
dtype: float64
```

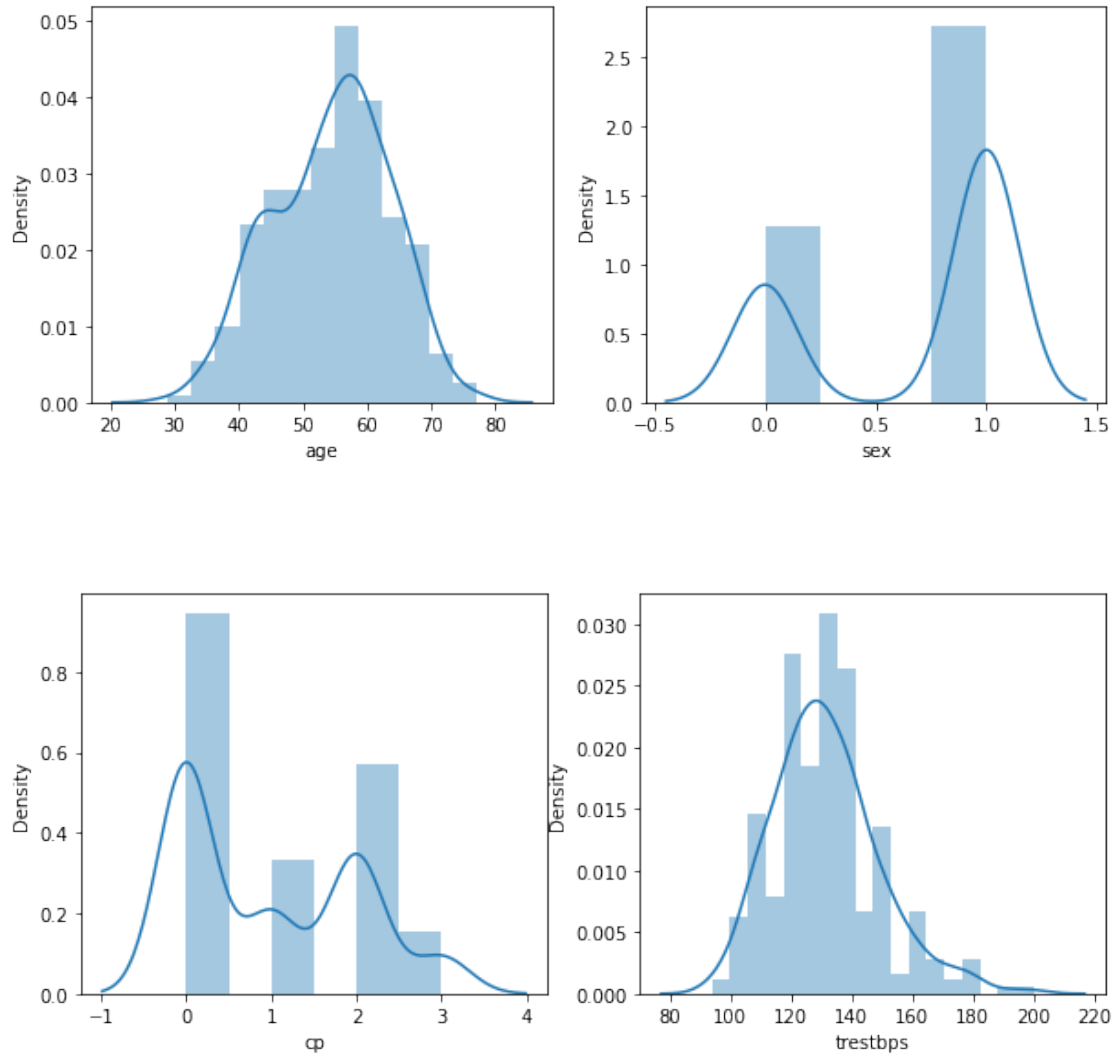
```
[13]: data.columns
```

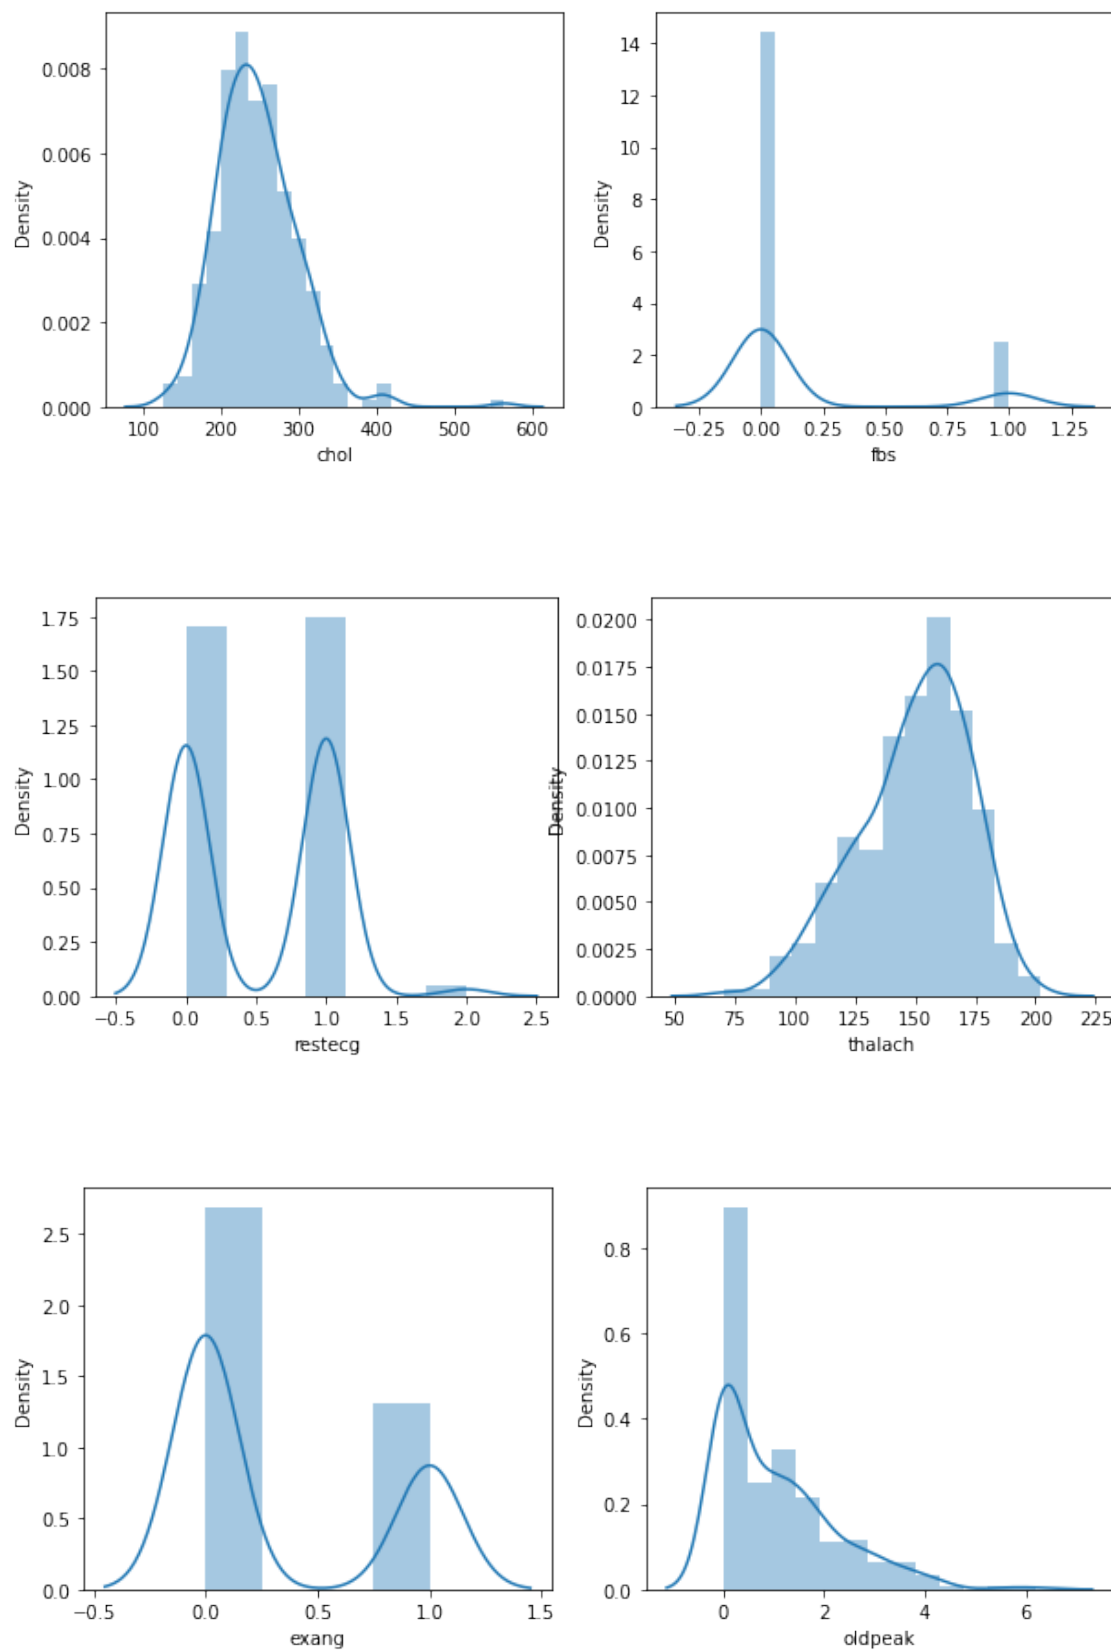
```
[13]: Index(['age', 'sex', 'cp', 'trestbps', 'chol', 'fbs', 'restecg', 'thalach',
          'exang', 'oldpeak', 'slope', 'ca', 'thal', 'target'],
          dtype='object')
```

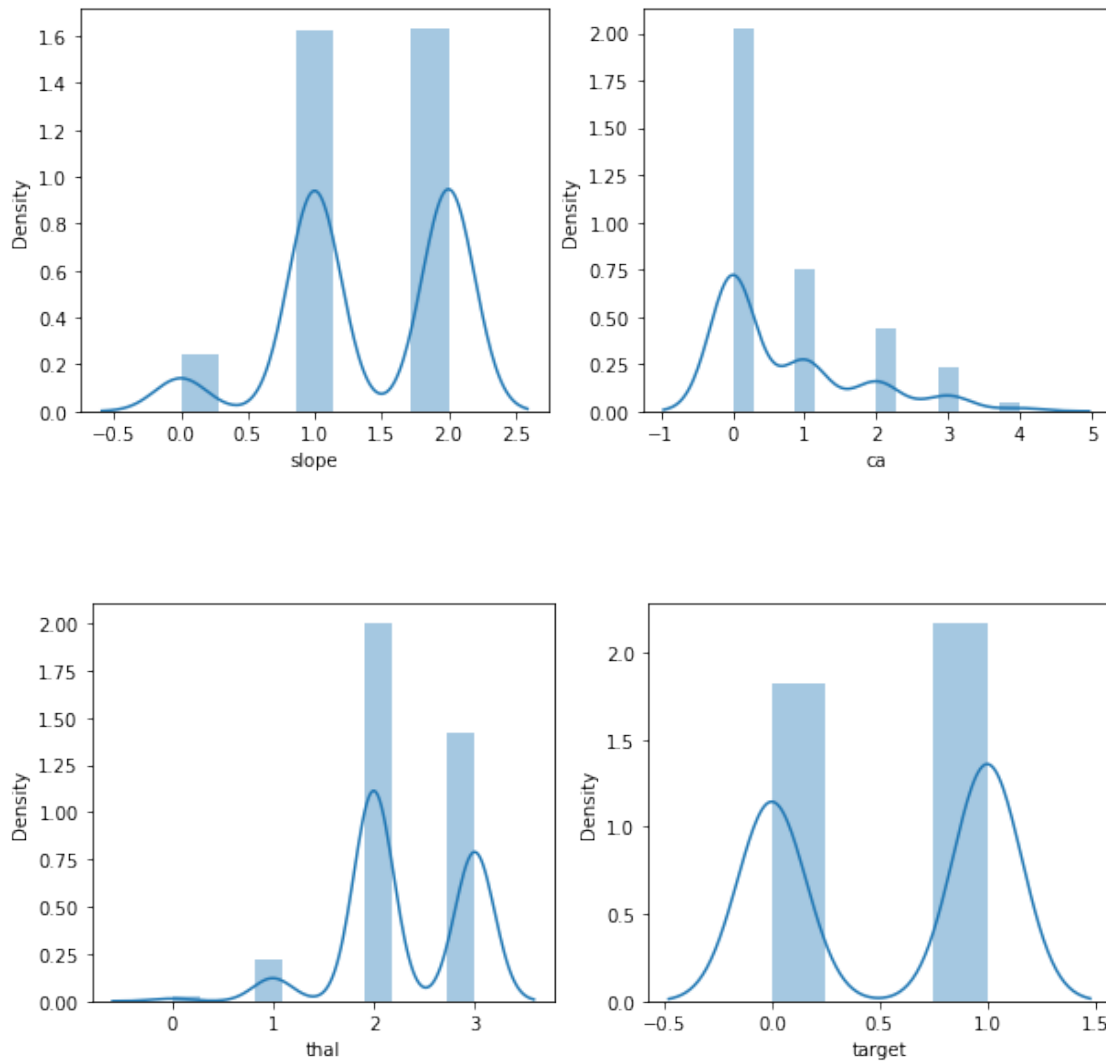
```
[14]: x= data.columns
```

```
[15]: for i in range(0,len(x)-1,2):
      plt.figure(figsize=(10,4))
```

```
plt.subplot(121)
sns.distplot(data[x[i]], kde= True)
plt.subplot(122)
sns.distplot(data[x[i+1]], kde= True)
```







4 Identify the data variables which are categorical and describe and explore these variables using the appropriate tools, such as count plot

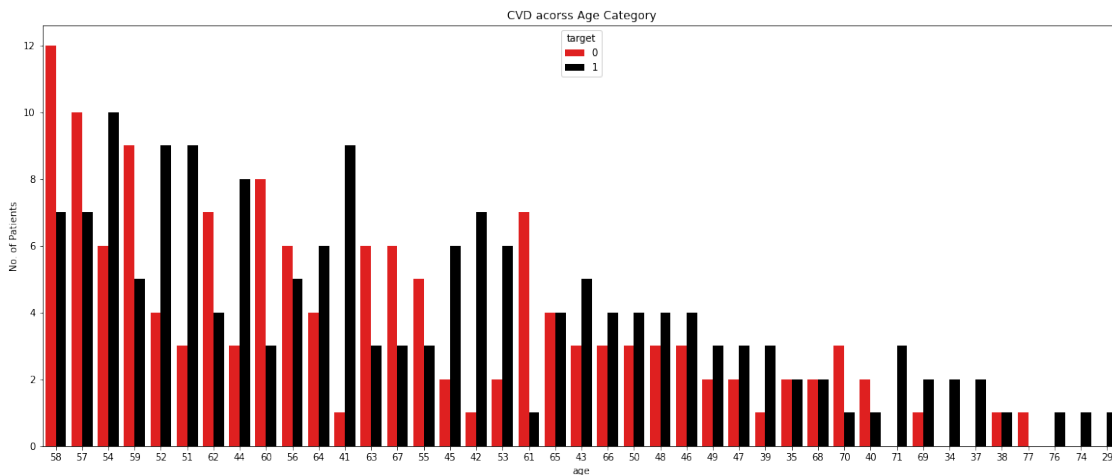
```
[16]: categorical_data = data.select_dtypes(exclude=[np.number])
```

```
[17]: categorical_data.sum()
```

```
[17]: Series([], dtype: float64)
```

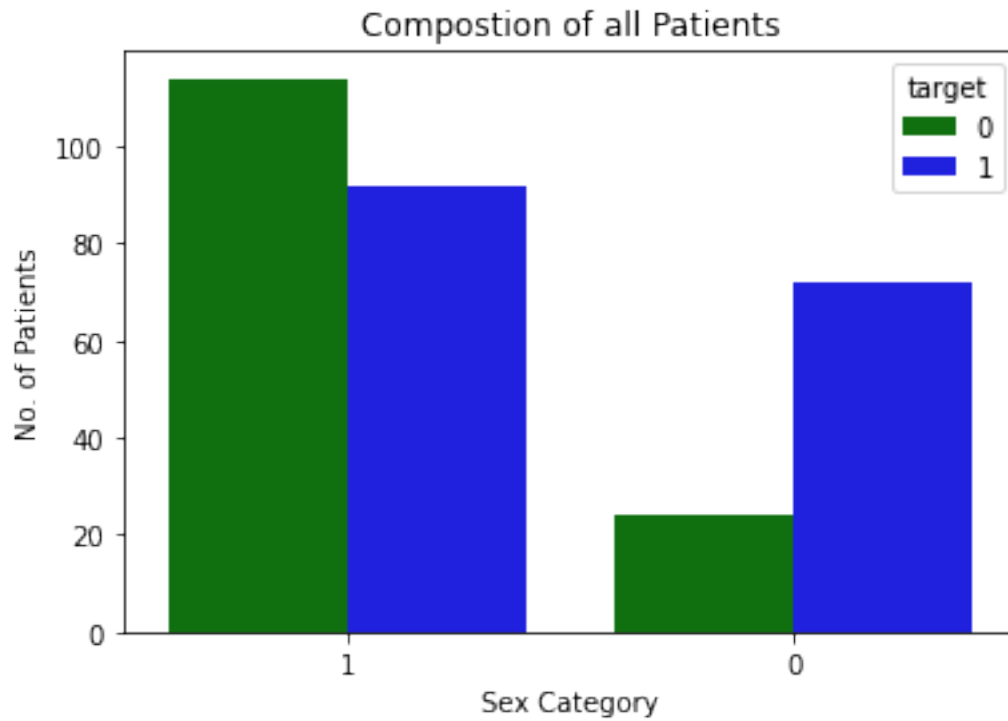
5 Study the occurrence of CVD across the Age category

```
[18]: plt.figure(figsize=(20,8))
plt.title("CVD across Age Category")
sns.
    ↳countplot(data=data,x="age",hue="target",palette=["red","black"],order=data['age']).
    ↳value_counts(ascending= False).index)
plt.ylabel("No. of Patients")
plt.show()
```




6 Study the composition of all patients with respect to the Sex category

```
[19]: ax=sns.countplot(data=data,x= data["sex"],hue="target",order=data["sex"].
    ↳value_counts(ascending= False).index, palette=["green","blue"]);
plt.title("Compostion of all Patients")
plt.xlabel("Sex Category")
plt.ylabel("No. of Patients")
plt.show()
```


7 Study if one can detect heart attacks based on anomalies in the resting blood pressure (trestbps) of a patient

```
[20]: plt.figure(figsize=(20,15))
sns.barplot(x="trestbps",y="target",data=data)
plt.show()
```



8 Describe the relationship between cholesterol levels and a target variable

```
[21]: sns.jointplot(x="chol",y="target",data=data,kind="hex")
plt.xlabel("Cholesterol")
plt.show()
```



```
[22]: data[data["sex"]==1]["target"]
```

```
[22]: 0      1
      1      1
      3      1
      5      1
      7      1
      ..
     295     0
     297     0
     299     0
     300     0
     301     0
```

Name: target, Length: 206, dtype: int64

```
[23]: from scipy.stats import pearsonr
data["chol"].corr(data["target"])
```

```
[23]: -0.08143720051844144
```

```
[24]: corr=pearsonr(data["chol"],data["target"])
```


```
[25]: corr
```

```
[25]: (-0.08143720051844144, 0.15803697464249133)
```

```
[26]: k=4
data.corr().nlargest(k,"chol")["chol"]
```

```
[26]: chol      1.000000
age       0.207216
trestbps  0.125256
thal      0.096810
Name: chol, dtype: float64
```

```
[27]: plt.figure(figsize=(20,12))
sns.heatmap(data.corr(),annot=True)
plt.show()
```




9 State what relationship exists between peak exercising and the occurrence of a heart attack

```
[28]: data[["oldpeak","target"]].corr()
```

```
[28]:      oldpeak    target  
oldpeak  1.000000 -0.429146  
target   -0.429146  1.000000
```


```
[29]: plt.figure(figsize=(10,6))  
sns.barplot(data=data,y="oldpeak",x="target")  
plt.show()
```



When peak exercising is less than 0.75 then the person has Occurance of Heart Attack

10 Check if thalassemia is a major cause of CVD

```
[30]: plt.figure(figsize=(20,12))  
sns.histplot(data)  
plt.show()
```

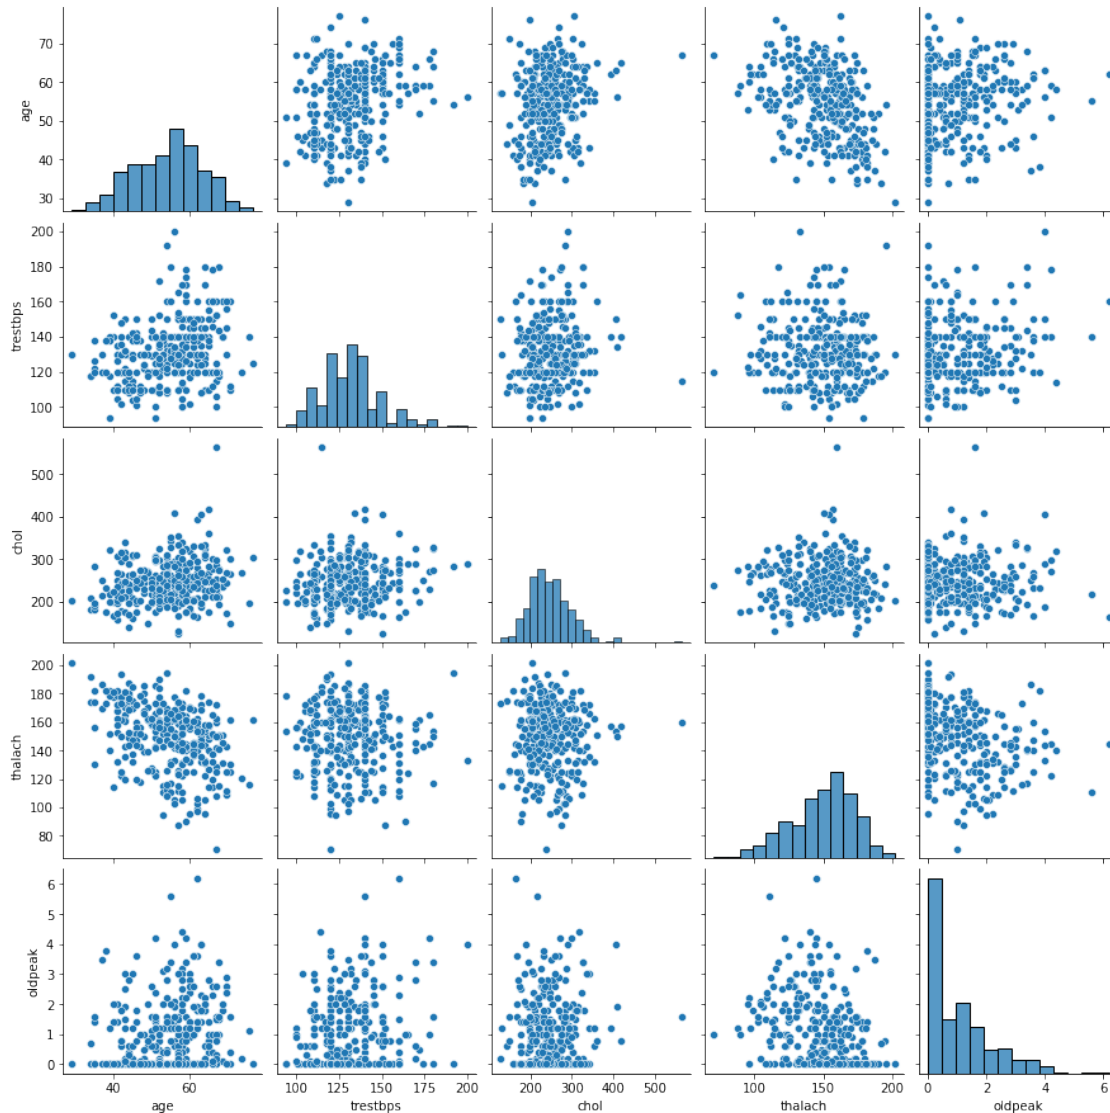


thalassemia is not the major cause of CVD

11 List how the other factors determine the occurrence of CVD

```
[31]: subdata=data[["age","trestbps","chol","thalach","oldpeak"]]  
sns.pairplot(subdata)
```

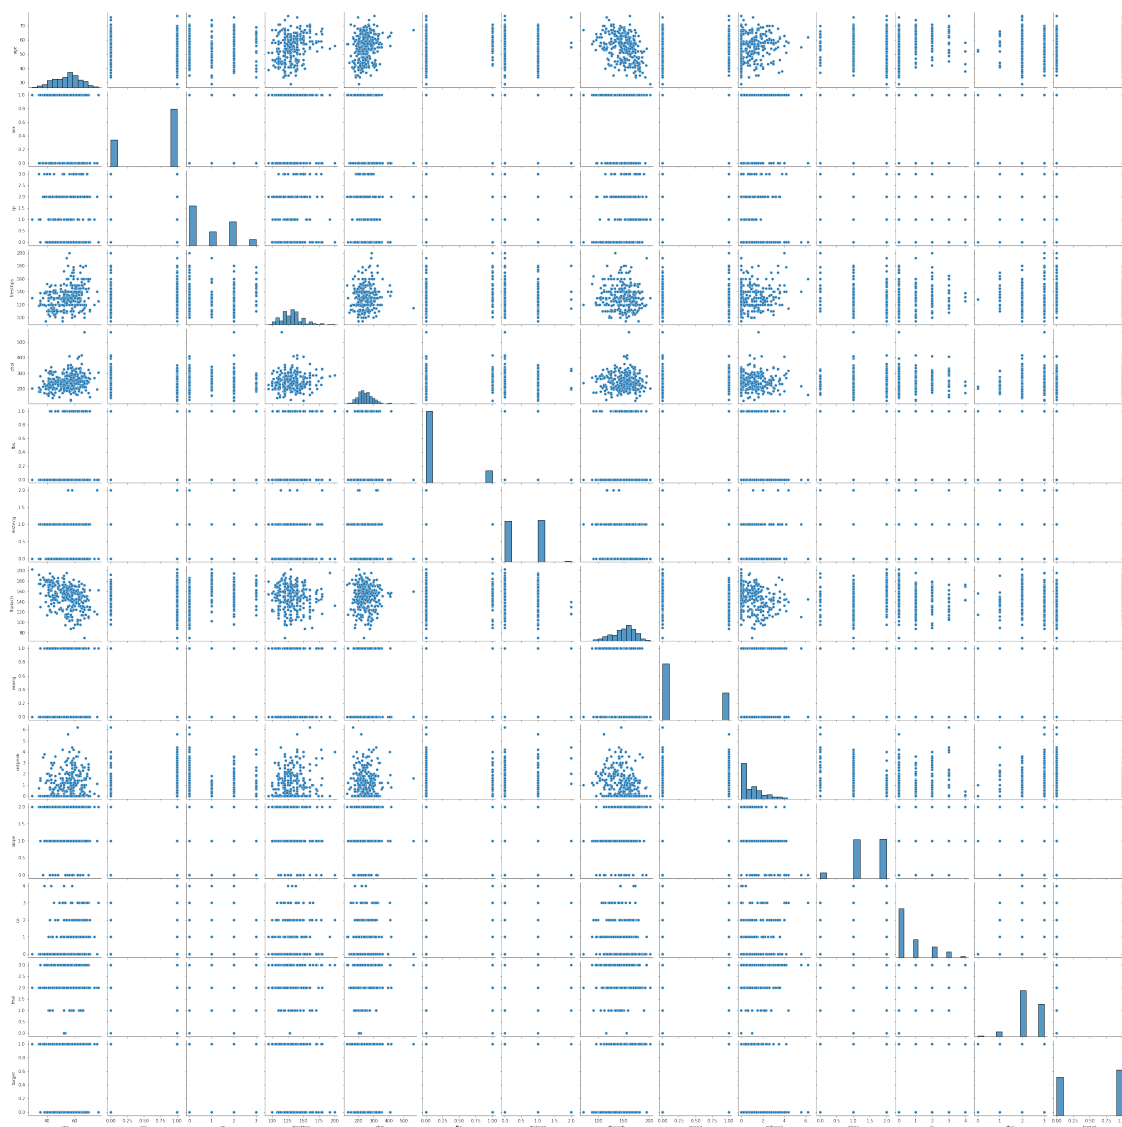
```
[31]: <seaborn.axisgrid.PairGrid at 0x7f5ec0b0e090>
```



12 Use a pair plot to understand the relationship between all the given variables

```
[32]: sns.pairplot(data)
```

```
[32]: <seaborn.axisgrid.PairGrid at 0x7f5ec1067ad0>
```



- 13 Build a baseline model to predict the risk of a heart attack using a logistic regression and random forest and explore the results while using correlation analysis and logistic regression (leveraging standard error and p-values from statsmodels) for feature selection

```
[33]: upper_limit_chol=data['chol'].mean()+3*data['chol'].std()
      lower_limit_chol=data['chol'].mean()-3*data['chol'].std()
```



```
[34]: data['chol']=np.where(  
    data['chol']>upper_limit_chol,  
    upper_limit_chol,  
    np.where(  
    data['chol']<lower_limit_chol,  
    lower_limit_chol,  
    data['chol']  
    )  
)
```

```
[35]: upper_limit_trestbps=data['trestbps'].mean()+3*data['trestbps'].std()  
lower_limit_trestbps=data['trestbps'].mean()-3*data['trestbps'].std()
```

```
[36]: data['trestbps']=np.where(  
    data['trestbps']>upper_limit_trestbps,  
    upper_limit_trestbps,  
    np.where(  
    data['trestbps']<lower_limit_trestbps,  
    lower_limit_trestbps,  
    data['trestbps']  
    )  
)
```

```
[37]: upper_limit_thalach=data['thalach'].mean()+3*data['thalach'].std()  
lower_limit_thalach=data['thalach'].mean()-3*data['thalach'].std()
```

```
[38]: data['thalach']=np.where(  
    data['thalach']>upper_limit_thalach,  
    upper_limit_thalach,  
    np.where(  
    data['thalach']<lower_limit_thalach,  
    lower_limit_thalach,  
    data['thalach']  
    )  
)
```

```
[39]: upper_limit_oldpeak=data['oldpeak'].mean()+3*data['oldpeak'].std()  
lower_limit_oldpeak=data['oldpeak'].mean()-3*data['oldpeak'].std()
```

```
[40]: data['oldpeak']=np.where(  
    data['oldpeak']>upper_limit_oldpeak,  
    upper_limit_oldpeak,  
    np.where(  
    data['oldpeak']<lower_limit_oldpeak,  
    lower_limit_oldpeak,  
    data['oldpeak']  
    )  
)
```

```

)

[41]: x=pd.DataFrame(data.iloc[:, :-1])
      y=pd.DataFrame(data.iloc[:, -1])

[42]: from sklearn.model_selection import train_test_split

[43]: from sklearn.preprocessing import StandardScaler

[44]: ss=StandardScaler()

[45]: ss.fit_transform(x)

[45]: array([[ 0.94979429,  0.68265615,  1.97647049, ..., -2.27118179,
              -0.71491124, -2.1479552 ],
             [-1.92854796,  0.68265615,  1.005911  , ..., -2.27118179,
              -0.71491124, -0.51399432],
             [-1.48572607, -1.46486632,  0.0353515 , ...,  0.97951442,
              -0.71491124, -0.51399432],
             ...,
             [ 1.50332164,  0.68265615, -0.93520799, ..., -0.64583368,
               1.27497996,  1.11996657],
             [ 0.28556146,  0.68265615, -0.93520799, ..., -0.64583368,
               0.28003436,  1.11996657],
             [ 0.28556146, -1.46486632,  0.0353515 , ..., -0.64583368,
               0.28003436, -0.51399432]])

[46]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=0)

[47]: from sklearn.linear_model import LogisticRegression

[48]: lr=LogisticRegression()

[49]: lr.fit(x_train,y_train)

[49]: LogisticRegression()

[50]: y_pred=lr.predict(x_test)
      y_pred

[50]: array([0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 1, 1, 0, 1, 1, 1, 1, 0,
              0, 0, 1, 0, 0, 0, 1, 1, 0, 0, 1, 1, 1, 0, 1, 1, 0, 0, 1, 1, 0, 0,
              0, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1])

[51]: from sklearn.metrics import accuracy_score
      accuracy_score(y_test,y_pred)

```

```
[51]: 0.8524590163934426
```

```
[52]: from sklearn.ensemble import RandomForestClassifier  
rfc=RandomForestClassifier()
```

```
[53]: rfc.fit(x_train,y_train)
```

```
[53]: RandomForestClassifier()
```

```
[54]: y_pred1=rfc.predict(x_test)  
y_pred1
```

```
[54]: array([0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 1, 1, 0, 1, 0, 1, 1, 0,  
        0, 0, 1, 0, 0, 0, 1, 1, 0, 0, 1, 1, 1, 0, 0, 1, 0, 0, 1, 0, 1, 0,  
        0, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1])
```

```
[55]: from sklearn.metrics import accuracy_score
```

```
[56]: accuracy_score(y_test,y_pred1)
```

```
[56]: 0.8852459016393442
```

```
[57]: from sklearn.metrics import r2_score  
r2_score(y_test,y_pred)
```

```
[57]: 0.4019607843137255
```

```
[58]: from scipy.stats import chi2_contingency  
#defining the table  
stat, p, dof, expected = chi2_contingency(data)  
# interpret p-value  
alpha = 0.05  
print("p value is " + str(p))  
if p <= alpha:  
    print('Dependent (reject H0)')  
else:  
    print('Independent (H0 holds true)')
```

```
p value is 3.125482768461386e-51  
Dependent (reject H0)
```

```
[ ]:
```